

Product Plots

By Hadley Wickham and Heike Hofmann

Jess Rees, Data Science Lab

August 12, 2018

Contents

Framework	1
Constraints:	1
Primitives:	1
Plots that fit into this framework:	2
Plot display	2
Extensions	2
Continuous data	2
Non-rectangular partitions	2
Using the package	3
Examples from the paper	3
Basic examples	4
More complex examples	8
Using gapminder data	19

Framework

Explained in detail in <http://vita.had.co.nz/papers/prodplots.pdf>

A new framework for visualising tables of counts, proportions and probabilities. Foundations:

- computation of area as a product of height and width
- statistical concept of generating a joint distribution from the product of conditional and marginal distributions

Focusses on area charts, where the area of a graphical element is proportional to the underlying count, proportion, or probability. Key development of the products plot framework is the inverse operation: the factorisation of high-dimensional data to products of low-dimensional plots.

Constraints:

- **area must be proportional to count** (or proportion, or probability)
- **partitions must be disjoint**; to see the complete area, each rectangle must be non-overlapping. This does not imply that the tiling must be space-filling
- **partitions must be rectangular**, because then many perceptual tasks only require comparing lengths, or positions along a common scale; these tasks are generally easier than comparing areas. Rectangles are also computationally simple, recursive (we can always tile a rectangle with smaller rectangles), and used in many existing visualisations.

Primitives:

- **bars (1d)**: *height is proportional to value*, while width equally divides space. Can be horizontally (hbar) or vertically (vbar) arranged.

- **spines (1d)**: *width is proportional to value*, and height fills the range. Space-filling. Can be horizontally (hspine) or vertically (vspine) arranged, or automatically by splitting the largest dimension (spine).
- **tiles (1d)**: *area is proportional to value*, with no restrictions on height or width other than trying to keep the aspect ratio of each rectangle close to 1.
- **fluct (2d)**: has *height and width proportional to the square root of the value*. Each rectangle is arranged on a regular grid formed by the levels of the two variables, allowing comparisons both vertically and horizontally.

Note: Bars and spines are indistinguishable when the underlying data is evenly distributed across the categories. Comparison is the easiest with bars (comparing positions on a common scale). But spines and tiles work better recursively, since they occupy the complete space.

Plots that fit into this framework:

- bar chart (1 hbar)
- column chart (1 vbar)
- spine plot (1 spine)
- fluctuation (1 fluct)
- stacked bar chart (1 hbar and 1 vspine)
- nested bar chart (2 hbars)
- equal bin size plot (1 fluct and 1 vspine)
- mosaic plot (alternating hspines and vspines)
- double decker plot (n-1 hspines and 1 vspine)
- treemap (n spines)
- squarified treemap (n tiles)
- generalised treemap (any plot ending with a tile)

Plot display

For labelling, use a combination of colour and axis labels. **Axis labelling is not very well supported in the package.** Tool tips may be very helpful.

Extensions

Continuous data

Continuous data can be binned to make them discrete, either into bins of equal width, or bins with an equal number of points. Leads to histograms (analogue to bar chart) and spinograms (analogue to spine chart)

Non-rectangular partitions

Radial plots can be seen as **polar transformations** of product plots. For example, a pie chart is an hspine drawn in polar coordinates with the x-coordinate mapped to angle and the y-coordinate to radius. Generally, the y axis (mapped to radius) must be square-root transformed to ensure that that counts stay proportional to areas.

- concentric pie chart (1 hspine)
- doughnut plot (1 hspine and 1 vspine)
- racetrack plot (1 vbar)
- infoslides (n vbars, using half of the polar plane)

Note that research suggests that **visualisations in polar coordinates are harder to read accurately than visualisations in Cartesian coordinates.**

Using the package

Available from <http://github.com/hadley/productplots>. Two main functions:

- **prodcalc**: computes the coordinates of each rectangle
- **prodplot**: displays the rectangles with ggplot2

Each graphical primitive is represented by a function: **hspine()**, **vspine()**, **spine()**, **hbar()**, **vbar()**, **tile()**, and **fluct()**.

```
knitr::opts_chunk$set(echo = TRUE)
#devtools::install_github("hadley/productplots")
library(productplots)
library(plyr)
library(tidyverse)
```

```
## -- Attaching packages ----- tid
## v ggplot2 3.0.0      v purrr  0.2.4
## v tibble  1.4.2      v dplyr  0.7.4
## v tidyr   0.8.0      v stringr 1.3.1
## v readr   1.1.1      v forcats 0.3.0

## -- Conflicts ----- tidyverse
## x dplyr::arrange()   masks plyr::arrange()
## x purrr::compact()  masks plyr::compact()
## x dplyr::count()     masks plyr::count()
## x dplyr::failwith()  masks plyr::failwith()
## x dplyr::filter()    masks stats::filter()
## x dplyr::id()        masks plyr::id()
## x dplyr::lag()       masks stats::lag()
## x dplyr::mutate()    masks plyr::mutate()
## x dplyr::rename()    masks plyr::rename()
## x dplyr::summarise() masks plyr::summarise()
## x dplyr::summarize() masks plyr::summarize()

library(ggplot2)
```

Examples from the paper

Using the happiness dataset:

```
getwd()

## [1] "C:/Users/root/Desktop"

d <- load("happy.rda") # downloaded from github; reference directory as needed
happy <- happy %>% filter(!is.na(happy))
head(happy)

##   id      happy year age  sex      marital      degree
## 1  1 not too happy 1972  23 female never married bachelor
## 2  2 not too happy 1972  70  male      married lt high school
```

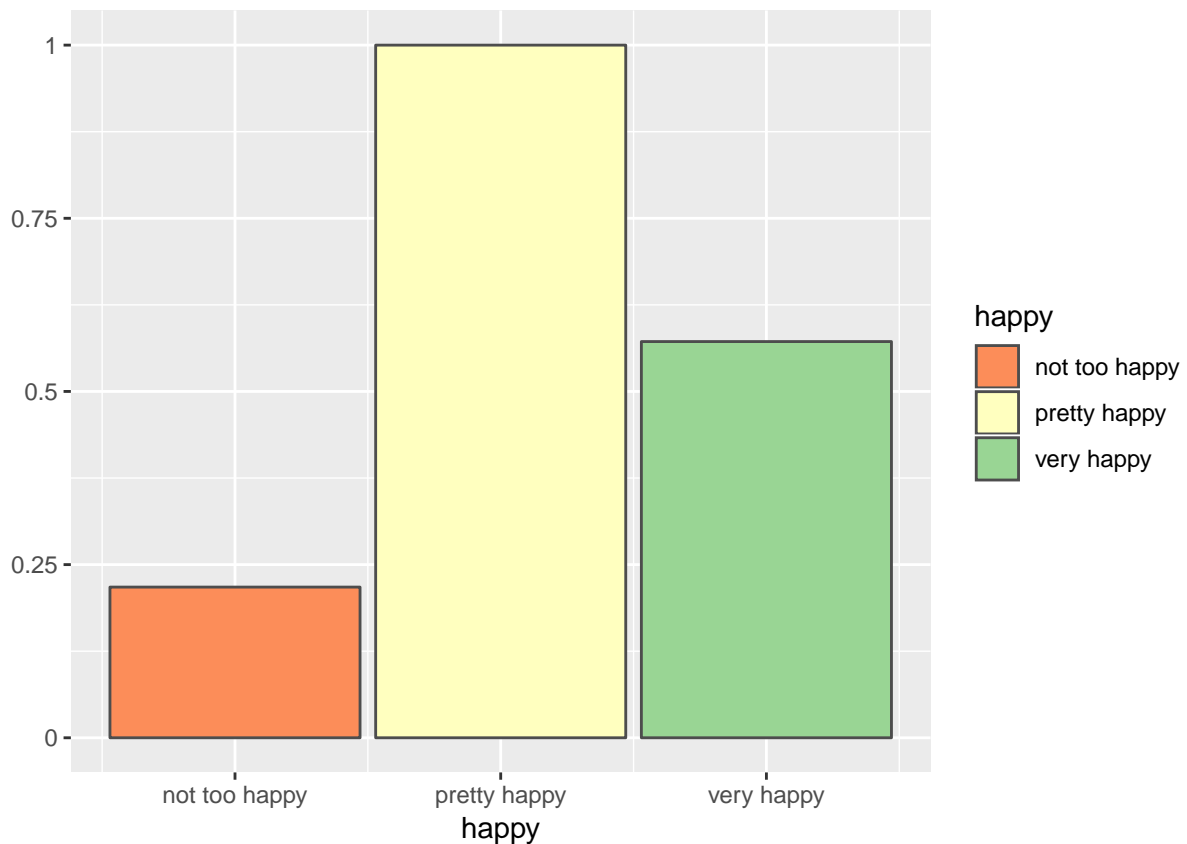
```
## 3 3 pretty happy 1972 48 female married high school
## 4 4 not too happy 1972 27 female married bachelor
## 5 5 pretty happy 1972 61 female married high school
## 6 6 pretty happy 1972 26 male never married high school
##      finrela health wtssall
## 1      average      good 0.4446
## 2 above average      fair 0.8893
## 3      average excellent 0.8893
## 4      average      good 0.8893
## 5 above average      good 0.8893
## 6 above average      good 0.4446
```

Basic examples

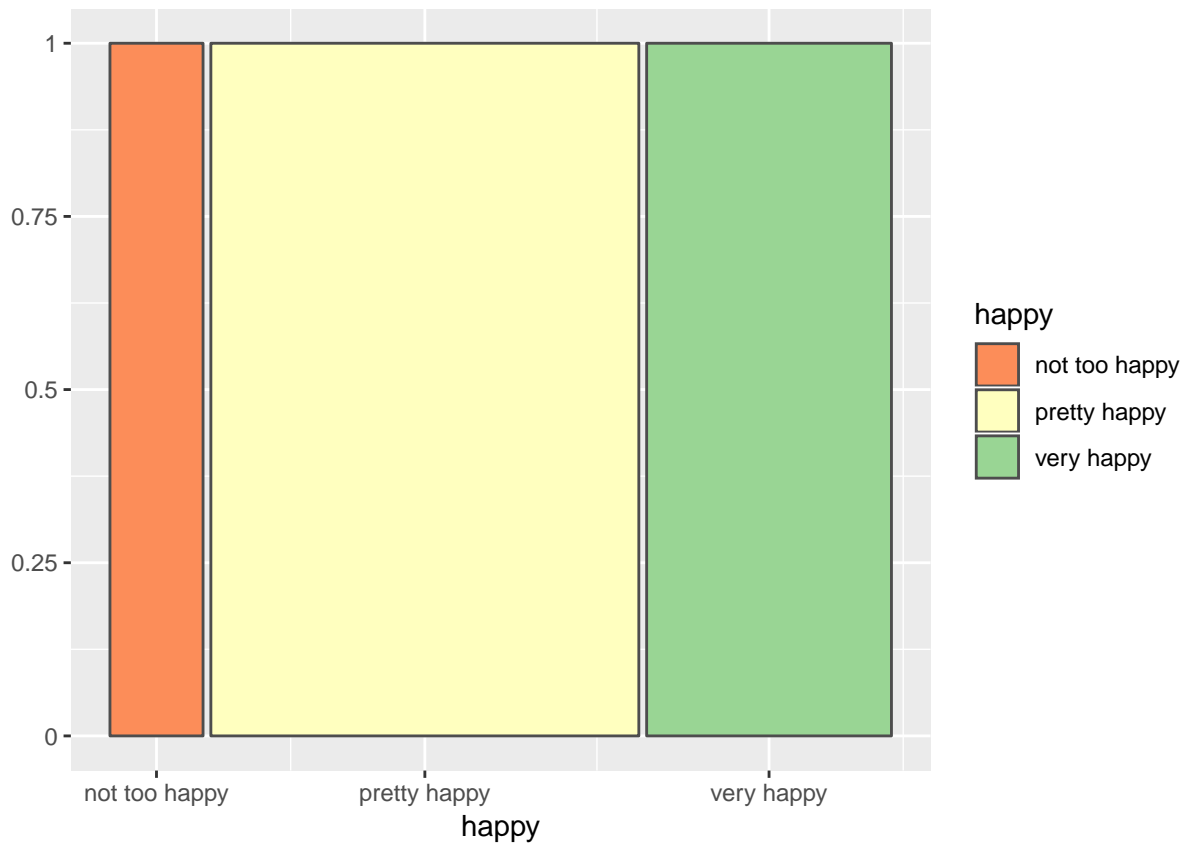
```
prodcalc(happy, ~ happy, "hbar")
```

```
##      happy .wt l r b t level
## 1 not too happy 5629 0.00 0.32 0 0.2175543 1
## 2 pretty happy 25874 0.34 0.66 0 1.0000000 1
## 3 very happy 14800 0.68 1.00 0 0.5720028 1
```


```
prodplot(happy, ~ happy, "hbar") + aes(fill=happy) +
  scale_fill_brewer(palette="Spectral")
```




```
prodplot(happy, ~ happy, "hspine") + aes(fill=happy) +
  scale_fill_brewer(palette="Spectral")
```



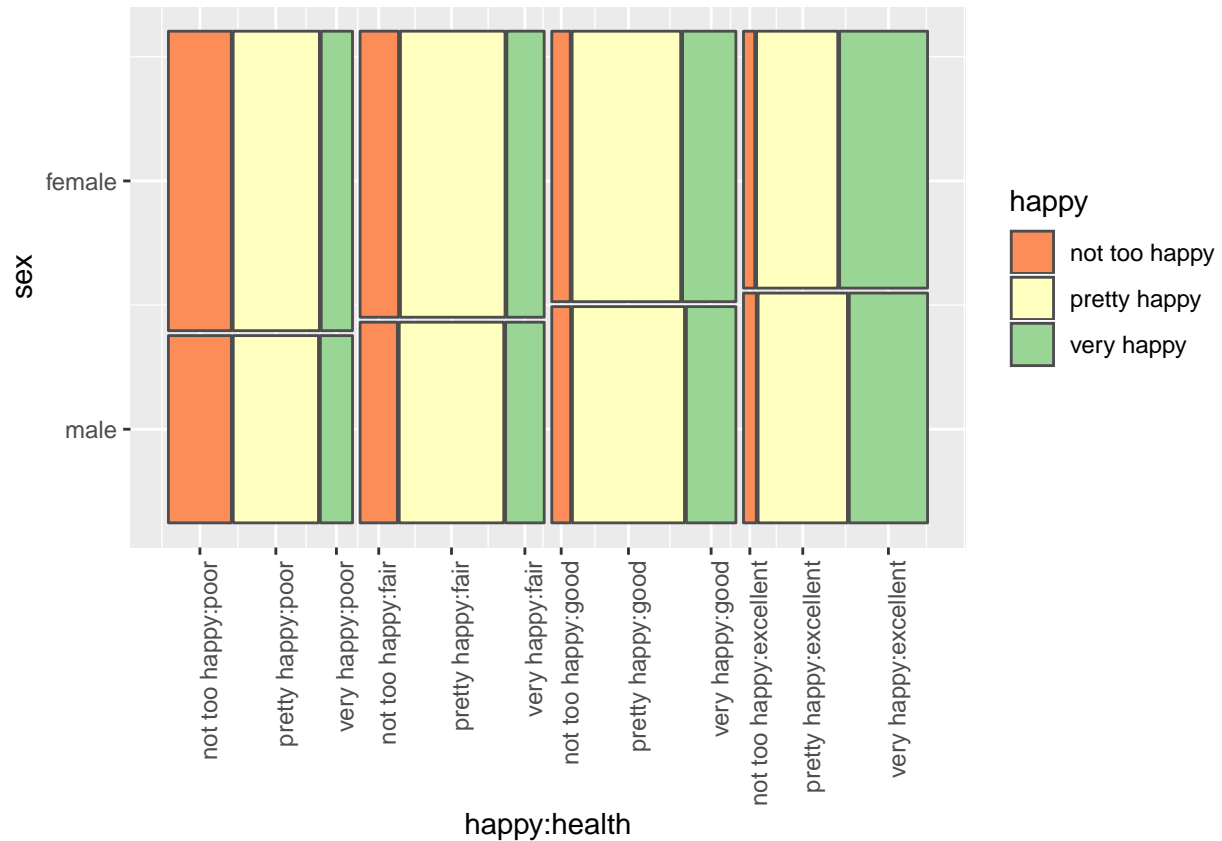
```
prodplot(happy, ~ sex + happy, c("vspine", "hbar")) + aes(fill=sex) +  
  scale_fill_brewer(palette="Spectral")
```



```
prodplot(happy, ~ sex + happy, stacked()) + aes(fill=happy) +  
  scale_fill_brewer(palette="Spectral")
```



```
prodplot(happy %>% filter(!is.na(health)), ~ happy + sex | health, mosaic("h")) +
  aes(fill=happy) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  scale_fill_brewer(palette="Spectral")
```



More complex examples


We'll use `fluct` and `spines` to see the relationship between happiness, health, and financial status.

`f(happy,health, finrela)`, partitioned with a `vspine` and `fluct`

$f(\text{happy}, \text{health}, \text{finrela}) = f(\text{happy}|\text{health}, \text{finrela}) \times f(\text{health}, \text{finrela})$

```
newhappy <- happy %>%
  mutate(finrela = as.factor(finrela)) %>%
  filter(!is.na(happy)) %>%
  filter(!is.na(finrela)) %>%
  filter(!is.na(health))

prodplot(newhappy, ~ happy + finrela + health, c("vspine", "fluct")) +
  aes(fill=happy) +
  scale_fill_brewer(palette="Spectral")
```

This plot displays raw proportions, showing that most people are in good health and average financial standing. However, it is difficult to see how happiness varies within these conditions because we must compare areas, not positions.

Health is on the x-axis, financial status on the y-axis. **But the labels are not right...**

```
# instead of calling prodplot (https://github.com/hadley/productplots/blob/master/R/plot.r), we go through
# this is all just what the function would do
levels = -1L
cascade=0
scale_max = T
na.rm = F
data = newhappy

formula = as.formula("~ happy + finrela + health") # same formula as above
vars <- parse_product_formula(formula)
p <- length(c(vars$cond, vars$marg))

divider <- c("vspine", "fluct") # same as above
if (is.function(divider)) divider <- divider(p)
div_names <- divider
if (is.character(divider)) divider <- llply(divider, match.fun)

# https://github.com/hadley/productplots/blob/master/R/calculate.r
res <- prodcalc(data, formula, divider, cascade, scale_max, na.rm = na.rm) # available from product plots

if (!(length(levels) == 1 && is.na(levels))) {
  levels[levels < 0] <- max(res$level) + 1 + levels[levels < 0]
}
```

```

res <- res[res$level %in% levels, ]
}

# here is where we change it; prodplot calls "draw", but we need a new draw function
df = list(data=res, formula=formula, divider=div_names)
alpha = 1
colour = "grey30"
subset = NULL
data <- df$data

# look at data
data

```

##		finrela	health	.wt	l	b	r
## 21	far	below average	poor	165	0.00	0.000000000	0.04476846
## 22	far	below average	poor	97	0.00	0.019726870	0.04476846
## 23	far	below average	poor	37	0.00	0.031471478	0.04476846
## 24		below average	poor	282	0.00	0.200000000	0.07286167
## 25		below average	poor	369	0.00	0.220922340	0.07286167
## 26		below average	poor	141	0.00	0.248119616	0.07286167
## 27		average	poor	190	0.00	0.400000000	0.06999955
## 28		average	poor	404	0.00	0.414824200	0.06999955
## 29		average	poor	137	0.00	0.445714398	0.06999955
## 30		above average	poor	31	0.00	0.600000000	0.02776422
## 31		above average	poor	61	0.00	0.606089780	0.02776422
## 32		above average	poor	23	0.00	0.617857946	0.02776422
## 33	far	above average	poor	13	0.00	0.800000000	0.01320150
## 34	far	above average	poor	8	0.00	0.805280599	0.01320150
## 35	far	above average	poor	5	0.00	0.808570819	0.01320150
## 36	far	below average	fair	186	0.25	0.000000000	0.30642653
## 37	far	below average	fair	231	0.25	0.017774239	0.30642653
## 38	far	below average	fair	58	0.25	0.039739484	0.30642653
## 39		below average	fair	549	0.25	0.200000000	0.36753726
## 40		below average	fair	1174	0.25	0.225486594	0.36753726
## 41		below average	fair	338	0.25	0.278917514	0.36753726
## 42		average	fair	515	0.25	0.400000000	0.39721071
## 43		average	fair	1946	0.25	0.419562405	0.39721071
## 44		average	fair	772	0.25	0.490209339	0.39721071
## 45		above average	fair	108	0.25	0.600000000	0.31701530
## 46		above average	fair	383	0.25	0.609005256	0.31701530
## 47		above average	fair	179	0.25	0.639575434	0.31701530
## 48	far	above average	fair	19	0.25	0.800000000	0.27196862
## 49	far	above average	fair	38	0.25	0.804720813	0.27196862
## 50	far	above average	fair	15	0.25	0.813986689	0.27196862
## 51	far	below average	good	165	0.50	0.000000000	0.56498406
## 52	far	below average	good	347	0.50	0.013863266	0.56498406
## 53	far	below average	good	118	0.50	0.042444699	0.56498406
## 54		below average	good	553	0.50	0.200000000	0.65538476
## 55		below average	good	2297	0.50	0.219945846	0.65538476
## 56		below average	good	752	0.50	0.298874745	0.65538476
## 57		average	good	704	0.50	0.400000000	0.73750000
## 58		average	good	5174	0.50	0.417477516	0.73750000
## 59		average	good	2537	0.50	0.533863411	0.73750000
## 60		above average	good	175	0.50	0.600000000	0.63670454

```

## 61      above average      good 1626 0.50 0.607820990 0.63670454
## 62      above average      good  987 0.50 0.671421356 0.63670454
## 63 far above average      good   17 0.50 0.800000000 0.53661437
## 64 far above average      good  107 0.50 0.802732896 0.53661437
## 65 far above average      good   76 0.50 0.818383342 0.53661437
## 66 far below average excellent   70 0.75 0.000000000 0.79980093
## 67 far below average excellent  182 0.75 0.007785096 0.79980093
## 68 far below average excellent  118 0.75 0.027388894 0.79980093
## 69      below average excellent  216 0.75 0.200000000 0.86138835
## 70      below average excellent  992 0.75 0.211081787 0.86138835
## 71      below average excellent  643 0.75 0.258774534 0.86138835
## 72          average excellent  314 0.75 0.400000000 0.94331208
## 73          average excellent 2631 0.75 0.410082603 0.94331208
## 74          average excellent 2630 0.75 0.483152974 0.94331208
## 75      above average excellent  112 0.75 0.600000000 0.88731609
## 76      above average excellent 1225 0.75 0.605384860 0.88731609
## 77      above average excellent 1476 0.75 0.653365142 0.88731609
## 78 far above average excellent   15 0.75 0.800000000 0.79653051
## 79 far above average excellent  124 0.75 0.802066358 0.79653051
## 80 far above average excellent  184 0.75 0.816443277 0.79653051
##          t level          happy
## 21 0.019368722      2 not too happy
## 22 0.031113331      2  pretty happy
## 23 0.035814767      2   very happy
## 24 0.220339447      2 not too happy
## 25 0.247536722      2  pretty happy
## 26 0.258289339      2   very happy
## 27 0.414264204      2 not too happy
## 28 0.445154401      2  pretty happy
## 29 0.455999639      2   very happy
## 30 0.605867666      2 not too happy
## 31 0.617635832      2  pretty happy
## 32 0.622211376      2   very happy
## 33 0.805174987      2 not too happy
## 34 0.808465207      2  pretty happy
## 35 0.810561199      2   very happy
## 36 0.017322827      2 not too happy
## 37 0.039288072      2  pretty happy
## 38 0.045141226      2   very happy
## 39 0.224546296      2 not too happy
## 40 0.277977216      2  pretty happy
## 41 0.294029805      2   very happy
## 42 0.418384719      2 not too happy
## 43 0.489031653      2  pretty happy
## 44 0.517768568      2   very happy
## 45 0.608469133      2 not too happy
## 46 0.639039312      2  pretty happy
## 47 0.653612239      2   very happy
## 48 0.804545064      2 not too happy
## 49 0.813810940      2  pretty happy
## 50 0.817574897      2   very happy
## 51 0.013343393      2 not too happy
## 52 0.041924826      2  pretty happy
## 53 0.051987246      2   very happy

```

```
## 54 0.218702768      2 not too happy
## 55 0.297631667      2  pretty happy
## 56 0.324307804      2   very happy
## 57 0.415577516      2 not too happy
## 58 0.531963411      2  pretty happy
## 59 0.590000000      2   very happy
## 60 0.606727354      2 not too happy
## 61 0.670327719      2  pretty happy
## 62 0.709363632      2   very happy
## 63 0.802439982      2 not too happy
## 64 0.818090427      2  pretty happy
## 65 0.829291495      2   very happy
## 66 0.007386689      2 not too happy
## 67 0.026990487      2  pretty happy
## 68 0.039840741      2   very happy
## 69 0.210190680      2 not too happy
## 70 0.257883427      2  pretty happy
## 71 0.289110678      2   very happy
## 72 0.408536107      2 not too happy
## 73 0.481606477      2  pretty happy
## 74 0.554649663      2   very happy
## 75 0.604286332      2 not too happy
## 76 0.652266613      2  pretty happy
## 77 0.709852870      2   very happy
## 78 0.801694114      2 not too happy
## 79 0.816071033      2  pretty happy
## 80 0.837224408      2   very happy
```


```
finrelabels = data %>%
  group_by(finrela) %>%
  filter(b == min(b)) %>% # want the label at the bottom of the block
  select(finrela, b) %>%
  distinct()
```

```
plot <- ggplot(data,
  ggplot2::aes_string(xmin = "l", xmax = "r", ymin = "b", ymax = "t")) +
  scale_y_product(df) + # from prodplots package
  scale_x_continuous(breaks = data$l, labels = data$health) + # put health labels at left of block
  scale_y_continuous(breaks = finrelabels$b, labels = finrelabels$finrela) # put finrela labels at bott
```

```
## Scale for 'y' is already present. Adding another scale for 'y', which
## will replace the existing scale.
```

```
# now we have to actually add the rectangles
levels <- split(data, data$level) # in this case, we only have 1
for (level in levels) {
  plot <- plot + geom_rect(data = level, colour = colour, alpha = alpha) #
}
```

```
plot +
  aes(fill=happy) + # as before
  scale_fill_brewer(palette="Spectral")
```



Now let's (kind of) generalise the code...

```
myprodplot <- function(data, str_formula, divider, forx, fory, forfill) {
  levels = -1L
  cascade=0
  scale_max = T
  na.rm = F
  alpha = 1
  colour = "grey30"
  subset = NULL

  formula = as.formula(str_formula)
  vars <- parse_product_formula(formula)
  p <- length(c(vars$cond, vars$marg))

  if (is.function(divider)) divider <- divider(p)
  div_names <- divider
  if (is.character(divider)) divider <- llply(divider, match.fun)

  res <- prodcalc(data, formula, divider, cascade, scale_max, na.rm = na.rm)

  if (!(length(levels) == 1 && is.na(levels))) {
    levels[levels < 0] <- max(res$level) + 1 + levels[levels < 0]
    res <- res[res$level %in% levels, ]
  }

  df = list(data=res, formula=formula, divider=div_names)
```

```

data <- df$data

colnum = which(colnames(data)==fory)
colnames(data)[colnum] <- "y"
ylabels = data %>%
  group_by(y) %>%
  filter(b == min(b)) %>%
  select(y, b) %>%
  distinct()

colnum = which(colnames(data)==forx)
colnames(data)[colnum] <- "x"
xlabels = data %>%
  group_by(x) %>%
  filter(l == min(l)) %>%
  select(x, l) %>%
  distinct()

plot <- ggplot(data,
  ggplot2::aes_string(xmin = "l", xmax = "r", ymin = "b", ymax = "t")) +
  scale_y_product(df) +
  scale_x_continuous(breaks = xlabels$l, labels = xlabels$x) +
  scale_y_continuous(breaks = ylabels$b, labels = ylabels$y)

levels <- split(data, data$level)
for (level in levels) {
  plot <- plot + geom_rect(data = level, colour = colour, alpha = alpha)
}


plot <- plot +
  ggplot2::aes_string(fill=forfill) +
  xlab(forx) +
  ylab(fory) +
  scale_fill_brewer(palette="Spectral")

return(plot)
}

p = myprodplot(newhappy, "~ happy + finrela + health", c("vspine", "fluct"), "health", "finrela", "happy")


## Scale for 'y' is already present. Adding another scale for 'y', which
## will replace the existing scale.
p

```



`f(happy | health, finrela), partitioned with a vspine and fluct`

```
prodplot(newhappy, ~ happy | finrela + health, c("vspine", "fluct")) +
  aes(fill=happy) +
  scale_fill_brewer(palette="Spectral")
```




Again, the labels are not what we need...

```
p = myprodplot(newhappy, "~ happy | finrela + health", c("vspine", "fluct"), "health", "finrela", "happy")
```

```
## Scale for 'y' is already present. Adding another scale for 'y', which
## will replace the existing scale.
```

```
p
```

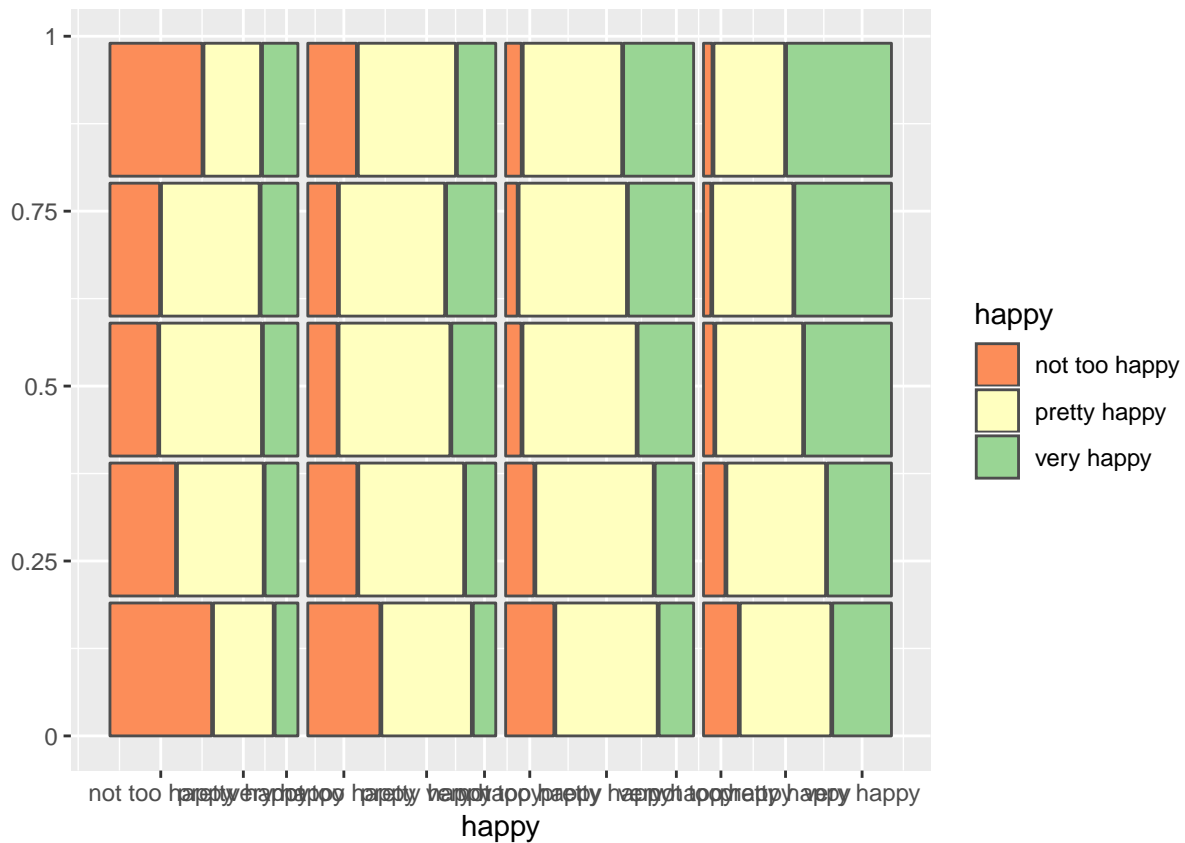



We can no longer see the joint distribution of health and financial status, but it is much easier to see the conditional distribution of happiness. Healthier and richer people are happier: maybe money does buy happiness?

Conditioning on financial status and health produces this plot (equal bin size plot) and makes it easier to see the conditional distribution of happiness given sex and health, because comparing positions along a common scale is an easier perceptual task. Depending on the comparison we are most interested in, we can make it easier to compare across wealth given health, or health given wealth, as in the next figure.

f(happy | health, finrela), partitioned with a hspine and fluct

```
prodplot(newhappy, ~ happy | finrela + health, c("hspine", "fluct")) +
  aes(fill=happy) +
  scale_fill_brewer(palette="Spectral")
```




Label issues...

```
p = myprodplot(newhappy, "~ happy | finrela + health", c("hspine", "fluct"), "health", "finrela", "happy")

## Scale for 'y' is already present. Adding another scale for 'y', which
## will replace the existing scale.

p
```



$f(\text{happy} \mid \text{health}, \text{finrela})$ partitioned with a `fluct` and `hspine`, emphasizing the relationship of happiness with finances, whereas the previous plot emphasizes the relationship with health.

Here we see that for a fixed income level, better health is correlated to increased happiness. The same is not true for a fixed level of health: rich people with poor health seem to be less happy than poorer people in poor health.


Using gapminder data

```
library(gapminder)
gm = gapminder

# first need to discretise
gm$yearcat = cut(gm$year, c(1950, 1960, 1970, 1980, 1990, 2000, 2010))
gm$lifeExpCat = cut(gm$lifeExp, c(20, 30, 40, 50, 60, 70, 80, 90))
gm$popCat = cut(gm$pop, c(0, 1000000, 10000000, 100000000, 1000000000, 10000000000))
gm$gdpPercapCat = cut(gm$gdpPercap, c(0, 2000, 4000, 6000, 8000, 10000, 50000, 100000, 150000))

p = myprodplot(gm, "~ lifeExpCat + yearcat + gdpPercapCat", c("vspine", "fluct"), "gdpPercapCat", "yearcat")

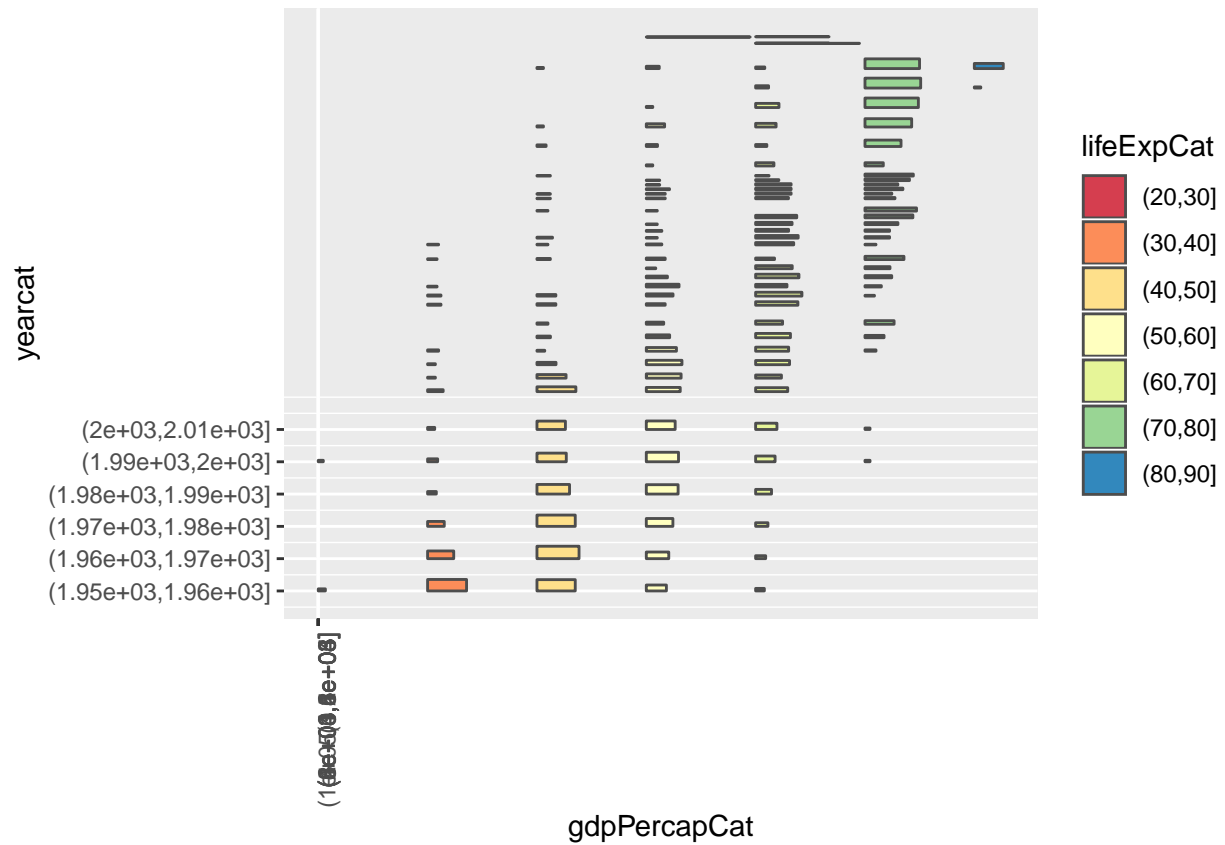
## Scale for 'y' is already present. Adding another scale for 'y', which
## will replace the existing scale.
p
```



```
r = myprodplot(gm, "~ lifeExpCat + yearcat + gdpPercapCat", c("fluct", "vspine"), "gdpPercapCat", "yearcat")

## Scale for 'y' is already present. Adding another scale for 'y', which
## will replace the existing scale.

r
```



```
q = myprodplot(gm, "~ gdpPercapCat | yearcat + continent", c("vspine", "fluct"), "continent", "yearcat"

## Scale for 'y' is already present. Adding another scale for 'y', which
## will replace the existing scale.

q
```

